

1. Project and Dataset Selection

For this project, we selected the Employer Data (Loan Approval and Risk Assessment) Kaggle dataset (<https://www.kaggle.com/datasets/gmudit/employer-data>). The dataset contains 20,000 rows and 36 features for demographic, employment, and financial information of loan applicants. It offers a classification problem (learning if a loan is approved) and regression problem (learning the applicant's risk score).

For this assignment, we applied the issue of classification in which the target variable is LoanApproved (0 = Not Approved, 1 = Approved). The predictors include factors such as Age, AnnualIncome, CreditScore, LoanAmount, DebtToIncomeRatio, etc., related to financial stability.

We intend to build and compare a number of tree-based models for loan approval outcome prediction and evaluate their performance based on metrics like Accuracy, Precision, Recall, F1-Score, and ROC-AUC.

2. Tree Model Building

2.1 Pre-Processing

The data was first loaded into a pandas DataFrame from a public URL.

Summary of data showed no missing values and no inconsistency in the 36 columns.

Categorical features such as EmploymentStatus, EducationLevel, MaritalStatus, HomeOwnershipStatus, and LoanPurpose were one-hot encoded.

The quantitative features were examined for distribution and correlation. Several key variables such as TotalDebtToIncomeRatio, AnnualIncome, and MonthlyIncome had right-skewed distributions since it would be true for most financial information. Features were normalized for comparison between models.

Correlation analysis indicated that DebtToIncomeRatio, CreditScore, and InterestRate were most correlated with loan approval, and they offered useful indications to learn from a model.

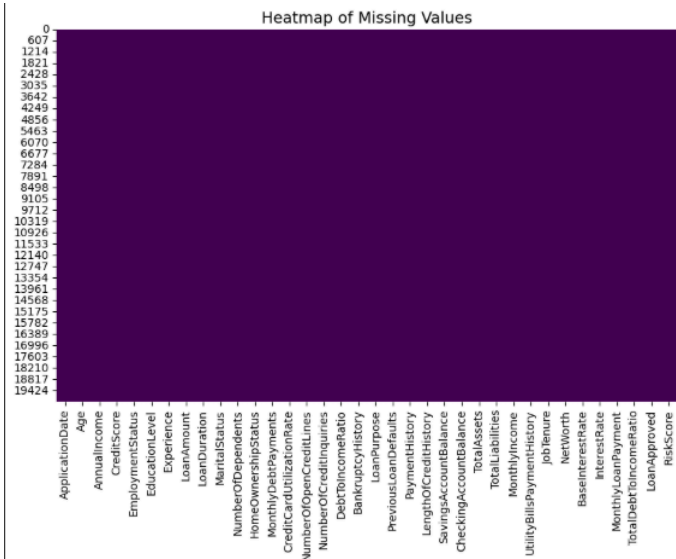
	ApplicationDate	Age	AnnualIncome	CreditScore	EmploymentStatus	EducationLevel	Experience	LoanAmount	LoanDuration	MaritalStatus	...	MonthlyIncome	UtilityBillPaymentHistory	JobTenure	NetWorth	BaseInterestRate	InterestRate	MonthlyLoanPayment	TotalDebtToIncomeRatio	LoanApproved	RiskScore
0	2016-01-01	45	39948	617	Employed	Master	22	15152	48	Married	...	3329.000000	0.724872	11	126028	0.199652	0.227598	419.805992	0.181877	0	49.0
1	2016-01-02	38	39709	628	Employed	Associate	15	26045	48	Single	...	3369.083333	0.935132	3	43609	0.287045	0.281877	794.054238	0.389852	0	52.0
2	2016-01-03	47	48724	570	Employed	Bachelor	26	17627	36	Married	...	3393.666667	0.872241	6	5205	0.217627	0.212548	666.486688	0.482157	0	52.0
3	2016-01-04	58	69084	545	Employed	High School	34	37898	96	Single	...	5757.000000	0.896155	5	99452	0.300388	0.300911	1047.506980	0.315898	0	54.0
4	2016-01-05	37	193264	594	Employed	Associate	17	9154	36	Married	...	8695.333333	0.941369	5	227019	0.197184	0.175990	336.179140	0.078210	1	36.0

5 rows × 36 columns

```

(20000, 36)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 36 columns):
 #   Column                Non-Null Count  Dtype  
---  --
 0   ApplicationDate        20000 non-null  object  
 1   Age                    20000 non-null  int64   
 2   AnnualIncome           20000 non-null  int64   
 3   CreditScore             20000 non-null  int64   
 4   EmploymentStatus       20000 non-null  object  
 5   EducationLevel         20000 non-null  object  
 6   Experience              20000 non-null  int64   
 7   LoanAmount             20000 non-null  int64   
 8   LoanDuration           20000 non-null  int64   
 9   MaritalStatus          20000 non-null  object  
10   NumberOfDependents     20000 non-null  int64   
11   HomeOwnershipStatus    20000 non-null  object  
12   MonthlyDebtPayments    20000 non-null  int64   
13   CreditCardUtilizationRate 20000 non-null  float64  
14   NumberOfOpenCreditLines 20000 non-null  int64   
15   NumberOfCreditInquiries 20000 non-null  int64   
16   DebtToIncomeRatio      20000 non-null  float64  
17   BankruptcyHistory      20000 non-null  int64   
18   LoanPurpose            20000 non-null  object  
19   PreviousLoanDefaults   20000 non-null  int64   
20   PaymentHistory         20000 non-null  int64   
21   LengthOfCreditHistory 20000 non-null  int64   
22   SavingsAccountBalance  20000 non-null  int64   
23   CheckingAccountBalance 20000 non-null  int64   
24   TotalAssets            20000 non-null  int64   
25   TotalLiabilities       20000 non-null  int64   
26   MonthlyIncome          20000 non-null  float64  
27   UtilityBillsPaymentHistory 20000 non-null  float64  
28   JobTenure              20000 non-null  int64   
29   NetWorth               20000 non-null  int64   
30   BaseInterestRate       20000 non-null  float64  
31   InterestRate           20000 non-null  float64  
32   MonthlyLoanPayment     20000 non-null  float64  
33   TotalDebtToIncomeRatio 20000 non-null  float64  
34   LoanApproved           20000 non-null  int64   
35   RiskScore              20000 non-null  float64  
dtypes: float64(9), int64(21), object(6)
memory usage: 5.5+ MB
None

```



2.2 Model Construction

Four tree-based classifiers were trained and tuned using GridSearchCV with cross-validation:

1. Decision Tree Classifier

```

Confusion Matrix:
[[4245  321]
 [ 349 1085]]

Classification Report:
              precision    recall  f1-score   support

     0       0.9240      0.9297      0.9269       4566
     1       0.7717      0.7566      0.7641       1434

   accuracy              0.8883        6000
  macro avg              0.8479      0.8432      0.8455        6000
 weighted avg              0.8876      0.8883      0.8880        6000

ROC AUC: 0.8431612958798615

```

2. Random Forest Classifier

```

Fitting 3 folds for each of 20 candidates, totalling 60 fits
RF Best Params: {'n_estimators': 300, 'min_samples_split': 5, 'min_samples_leaf': 2, 'max_features': None, 'max_depth': None, 'criterion': 'gini', 'bootstrap': True}

Confusion Matrix:
[[4370 186]
 [ 210 1224]]

Classification Report:

```

	precision	recall	f1-score	support
0	0.9541	0.9571	0.9556	4566
1	0.8620	0.8536	0.8577	1434
accuracy			0.9323	6000
macro avg	0.9081	0.9053	0.9067	6000
weighted avg	0.9321	0.9323	0.9322	6000

```

ROC AUC: 0.9785593187994266

```

3. AdaBoost Classifier

```

Fitting 3 folds for each of 12 candidates, totalling 36 fits
Ada Best Params: {'learning_rate': 0.5, 'n_estimators': 300}

Confusion Matrix:
[[4454 112]
 [ 190 1244]]

Classification Report:

```

	precision	recall	f1-score	support
0	0.9591	0.9755	0.9672	4566
1	0.9174	0.8675	0.8918	1434
accuracy			0.9497	6000
macro avg	0.9382	0.9215	0.9295	6000
weighted avg	0.9491	0.9497	0.9492	6000

```

ROC AUC: 0.9888642388010099

```

4. XGBoost Classifier

```

Fitting 3 folds for each of 20 candidates, totalling 60 fits
XGB Best Params: {'subsample': 0.9, 'min_child_weight': 5, 'max_depth': 3, 'learning_rate': 0.1, 'gamma': 1.0, 'colsample_bytree': 1.0}

Confusion Matrix:
[[4351 215]
 [  65 1369]]

Classification Report:

```

	precision	recall	f1-score	support
0	0.9853	0.9529	0.9688	4566
1	0.8643	0.9547	0.9072	1434
accuracy			0.9533	6000
macro avg	0.9248	0.9538	0.9380	6000
weighted avg	0.9564	0.9533	0.9541	6000

```

ROC AUC: 0.9926133736043071

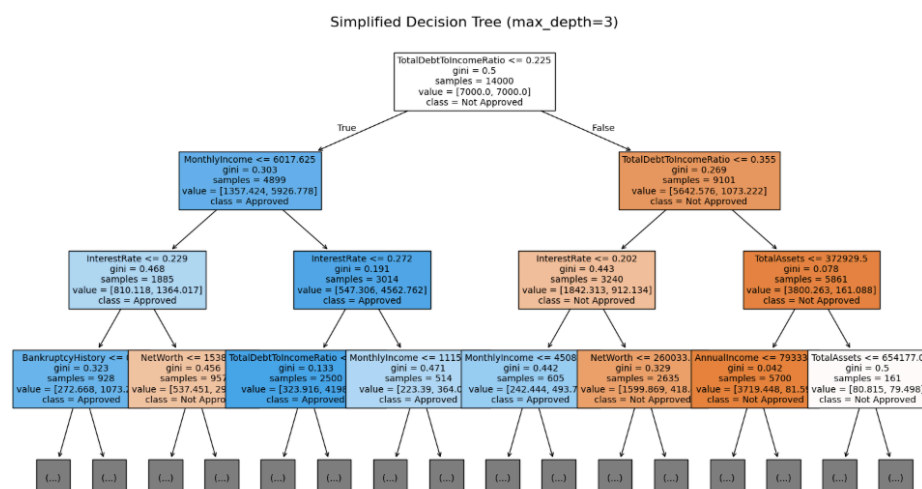
```

All models were evaluated on a train-test split (70-30) with balanced class weighting to address the slight imbalance in approval outcomes (76% Not Approved, 24% Approved). Hyperparameters such as max_depth, n_estimators, learning_rate, and min_samples_leaf were optimized for best performance.

2.3 Tree Visualization

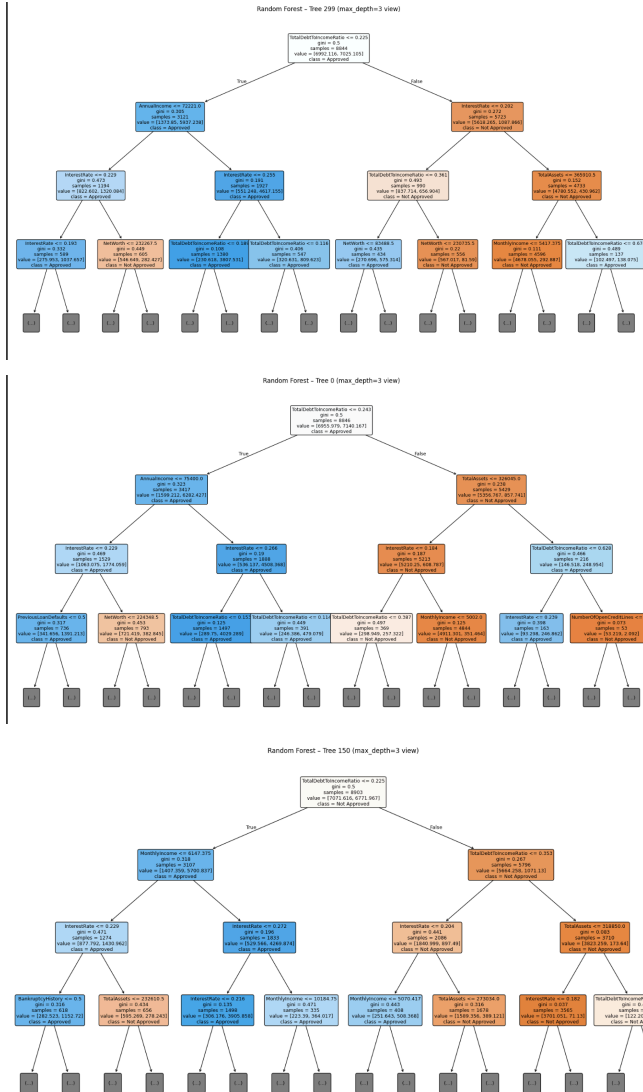
To understand how each model makes its predictions, decision trees were visualized for all four classifiers — Decision Tree, Random Forest, AdaBoost, and XGBoost.

The Decision Tree Classifier produced the most interpretable single model. The root node split was on TotalDebtToIncomeRatio, meaning that candidates with lower debt relative to income were more likely to be approved. Splits followed on MonthlyIncome, InterestRate, and NetWorth, reflecting significant lending reasoning, where higher income and lower interest rates increase approval probability. The lone tree readily shows how the model separates approved from not-approved candidates.



For the Random Forest Classifier, three representative trees (Tree 0, Tree 150, and Tree 299) were plotted, each limited to a maximum depth of 3 for readability.

- All trees began with TotalDebtToIncomeRatio as the dominant root split, confirming its importance across the ensemble.
- Secondary splits frequently appeared on AnnualIncome, MonthlyIncome, and TotalAssets, identifying high-earning and asset-rich applicants as lower-risk.
- Additional splits on InterestRate, NetWorth, and PreviousLoanDefaults refined subgroups of borderline applicants.
These trees demonstrate how Random Forest combines many shallow, diverse learners to reduce variance and achieve stable accuracy.



For the XGBoost Classifier, three trees (Tree 0, Tree 10, and Tree 399) were visualized to highlight the gradient-boosting process.

- Tree 0 began with `TotalDebtToIncomeRatio`, followed by `MonthlyIncome`, `InterestRate`, and `TotalAssets`, reflecting the model's foundational decision pattern.
- Tree 10 refined these patterns with new splits on `TotalAssets` and `AnnualIncome`, capturing subtle income-to-debt relationships.
- Tree 399 focused on `PreviousLoanDefaults`, `CreditScore`, and `SavingsAccountBalance`, showing how later iterations fine-tune edge cases to optimize performance. Together, these trees illustrate XGBoost's ability to learn from residual errors and incrementally build a highly accurate ensemble.



The AdaBoost Classifier visualizations (Base Estimators 0, 150, and 299) illustrate the model's sequential learning behavior.

- Base Estimator 0 relied almost entirely on `TotalDebtToIncomeRatio`, drawing a single clear boundary between approved and not-approved applicants.
- Base Estimator 150 introduced `TotalAssets` as an additional criterion, suggesting that asset value becomes a key correction factor as boosting continues.
- Base Estimator 299 emphasized `NetWorth`, fine-tuning the model's focus on applicants with substantial assets.
These incremental changes reflect AdaBoost's mechanism — each successive tree corrects the mistakes of earlier ones, gradually improving predictive accuracy through weighted re-training.


Across all models, the most consistent and influential predictors were `TotalDebtToIncomeRatio`, `Income` (Monthly and Annual), `InterestRate`, and `TotalAssets`. The recurrence of these variables across different algorithms confirms their strong predictive value in determining loan-approval outcomes.

2.4 Result Analysis

Model performance was evaluated on the test set using Accuracy, Precision, Recall, F1-Score, and ROC-AUC metrics.

The XGBoost Classifier achieved the best results with the highest AUC (0.993) and a well-balanced precision-recall tradeoff. Both XGBoost and AdaBoost outperformed the Random Forest and Decision Tree, demonstrating that boosting methods handle bias and variance more effectively.

Feature importance plots further highlighted TotalDebtToIncomeRatio, AnnualIncome, and LoanDuration as the most significant predictors of loan approval. This aligns closely with the visualized trees, where these same attributes consistently appeared near the top splits.



	Accuracy	Precision	Recall	F1	ROC AUC
XGBoost	0.9533	0.8643	0.9547	0.9072	0.9926
AdaBoost	0.9497	0.9174	0.8675	0.8918	0.9889
Random Forest	0.9323	0.8620	0.8536	0.8577	0.9786
Decision Tree	0.8908	0.7250	0.8752	0.7930	0.9321

3. Conclusion

For all experiments, XGBoost was the top performer with enhanced generalization power and most stable performance across metrics. The results suggest that income-based characteristics and debt ratio are key drivers of the probability of loan approval. The project demonstrated the end-to-end pipeline of:

- Data cleaning and transformation,
- Model tuning with GridSearchCV,
- Comparative evaluation of ensemble methods, and
- Visual interpretation through ROC curves and decision trees.

Future improvements could include using SMOTE for better class balance, SHAP analysis for model explainability, and exploring the regression aspect of predicting RiskScore.